

# Problem Statement

The objective of this case study is to develop a predictive model for LoanTap, an online platform specializing in customized loan products for millennials. The focus is on the underwriting process for Personal Loans. Given a set of attributes for an individual, the goal is to determine whether a credit line should be extended and, if so, to provide recommendations on the repayment terms. The model should accurately predict the likelihood of loan repayment versus default, helping LoanTap make informed lending decisions and minimize the risk of non-performing assets (NPAs).

## Imports and Dataset

### Imports

#### Imports for EDA

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

#### Encoding Imports

```
In [ ]: from sklearn.impute import SimpleImputer
```

#### Imports for Logistic Regression

```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
```

#### VIF import

```
In [ ]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

## Dataset

```
In [ ]: df = pd.read_csv("loan_tap.csv")
df.head()
```

```
Out [ ]: 
```

|   | loan_amnt | term      | int_rate | installment | grade | sub_grade | emp_title | emp_length | home_own |
|---|-----------|-----------|----------|-------------|-------|-----------|-----------|------------|----------|
| 0 | 10000.0   | 36 months | 11.44    | 329.48      | B     | B4        | Marketing | 10+ years  |          |

|   | loan_amnt | term      | int_rate | installment | grade | sub_grade | emp_title               | emp_length | home_own |
|---|-----------|-----------|----------|-------------|-------|-----------|-------------------------|------------|----------|
| 1 | 8000.0    | 36 months | 11.99    | 265.68      | B     | B5        | Credit analyst          | 4 years    | MOR      |
| 2 | 15600.0   | 36 months | 10.49    | 506.97      | B     | B3        | Statistician            | < 1 year   |          |
| 3 | 7200.0    | 36 months | 6.49     | 220.65      | A     | A2        | Client Advocate         | 6 years    |          |
| 4 | 24375.0   | 60 months | 17.27    | 609.33      | C     | C5        | Destiny Management Inc. | 9 years    | MOR      |

5 rows × 27 columns

# Data Preprocessing and EDA

## Basic Checks

```
In [ ]: df.shape
```

```
Out[ ]: (396030, 27)
```

```
In [ ]: df.duplicated().sum()
```

```
Out[ ]: 0
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            396030 non-null float64
1   term                                 396030 non-null object
2   int_rate                             396030 non-null float64
3   installment                           396030 non-null float64
4   grade                                 396030 non-null object
5   sub_grade                            396030 non-null object
6   emp_title                            373103 non-null object
7   emp_length                           377729 non-null object
8   home_ownership                       396030 non-null object
9   annual_inc                           396030 non-null float64
10  verification_status                  396030 non-null object
11  issue_d                              396030 non-null object
12  loan_status                           396030 non-null object
13  purpose                              396030 non-null object
14  title                                394275 non-null object
15  dti                                  396030 non-null float64
16  earliest_cr_line                     396030 non-null object
17  open_acc                             396030 non-null float64
18  pub_rec                              396030 non-null float64
```

```

19  revol_bal          396030 non-null  float64
20  revol_util         395754 non-null  float64
21  total_acc          396030 non-null  float64
22  initial_list_status 396030 non-null  object
23  application_type    396030 non-null  object
24  mort_acc            358235 non-null  float64
25  pub_rec_bankruptcies 395495 non-null  float64
26  address             396030 non-null  object

```

dtypes: float64(12), object(15)

memory usage: 81.6+ MB

- Total Number of rows: 396030
- Total Number of Columns: 27
- No duplicates.
- There are some missing values.

## Missing values Treatment

```

In [ ]: missing = df.isna().sum() * 100 / len(df)
missing = missing[(missing > 0)]

print('Columns with missing values')
missing

```

Columns with missing values

```

Out[ ]: emp_title      5.789208
emp_length  4.621115
title       0.443148
revol_util  0.069692
mort_acc    9.543469
pub_rec_bankruptcies 0.135091
dtype: float64

```

Removing rows which contains less than 3% missing values.

```

In [ ]: df.dropna(subset=(missing[missing <= 3]).index, inplace=True)

```

```

In [ ]: missing = df.isna().sum() * 100 / len(df)
missing = missing[(missing > 0)]

print('Columns with % missing values')
missing

```

Columns with % missing values

```

Out[ ]: emp_title      5.761122
emp_length  4.594055
mort_acc    9.453192
dtype: float64

```

Filling emp\_title missing values with 'Unknown'

```

In [ ]: df['emp_title'].fillna(value= 'Unknown', inplace=True)

```

Filling emp\_length missing values with most\_frequent value

```

In [ ]: emp_length_imputer = SimpleImputer(strategy="most_frequent")
df['emp_length'] = emp_length_imputer.fit_transform(df[['emp_length']]).flatten()

```

Filling mort\_acc missing values with

```
In [ ]: mort_acc_imputer = SimpleImputer(strategy="median")
df['mort_acc'] = mort_acc_imputer.fit_transform(df[['mort_acc']]).flatten()
```

```
In [ ]: missing = df.isna().sum() * 100 / len(df)
missing = missing[(missing > 0)]

print('Columns with % missing values')
missing
```

Columns with % missing values

```
Out[ ]: Series([], dtype: float64)
```

**Missing values has been treated.**

## Other Columns Treatment

### Change dtype of date/time columns

```
In [ ]: df['issue_d'][:5]
```

```
Out[ ]: 0    Jan-2015
1    Jan-2015
2    Jan-2015
3    Nov-2014
4    Apr-2013
Name: issue_d, dtype: object
```

```
In [ ]: df['earliest_cr_line'][:5]
```

```
Out[ ]: 0    Jun-1990
1    Jul-2004
2    Aug-2007
3    Sep-2006
4    Mar-1999
Name: earliest_cr_line, dtype: object
```

```
In [ ]: # Note below steps will add day in date value, so later on if we used these column,

df['issue_d'] = pd.to_datetime(df['issue_d'])
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'])
```

## Feature Engineering

### Flagging of pub\_rec, mort\_acc and pub\_rec\_bankruptcies

```
In [ ]: def pub_rec(n: float) -> int:
    return int(n != 0)

def mort_acc(n: float) -> int:
    return int(n != 0)

def pub_rec_bankruptcies(n: float) -> int:
    return int(n != 0)
```

```
In [ ]: df['pub_rec'] = df['pub_rec'].apply(pub_rec)
df['mort_acc'] = df['mort_acc'].apply(mort_acc)
df['pub_rec_bankruptcies'] = df['pub_rec_bankruptcies'].apply(pub_rec_bankruptcies)
```

## Extracting zip\_codes and states from states

```
In [ ]: # utility function to extract states
def find_state(s: str) -> str:
    ind = s.find(',')
    output = ""

    # finding abbreviation of state
    if ind != -1:
        output = s[ind+2 : ind+4]

    if output == "Bo" or output == "":
        output = "OTHER"

    return output

def find_zip_code(s: str) -> str:
    return s[-5:]
```

```
In [ ]: df['state'] = df['address'].apply(find_state)
df['zip_code'] = df['address'].apply(find_zip_code)

df.drop('address', axis=1, inplace=True) # dropping address column
```

## Date-time components

```
In [ ]: df['issue_d_month'] = df['issue_d'].dt.month
df['issue_d_year'] = df['issue_d'].dt.year

df['earliest_cr_line_month'] = df['earliest_cr_line'].dt.month
df['earliest_cr_line_year'] = df['earliest_cr_line'].dt.year
```

```
In [ ]: df.drop(['issue_d', 'earliest_cr_line'], axis=1, inplace=True) # dropping issue_d and earliest_cr_line
```

## Analysis and Visualisations

### value\_counts()

```
In [ ]: df['loan_status'].value_counts(normalize=True)
```

```
Out[ ]: Fully Paid      0.80381
Charged Off    0.19619
Name: loan_status, dtype: float64
```

```
In [ ]: df['home_ownership'].value_counts()
```

```
Out[ ]: MORTGAGE    197110
RENT              158770
OWN               37443
```

```
OTHER      110
NONE       29
ANY        3
Name: home_ownership, dtype: int64
```

combining none and any to other

```
In [ ]: df.loc[df['home_ownership'].isin(['NONE', 'ANY']), 'home_ownership'] = "OTHER"
df['home_ownership'].value_counts()
```

```
Out[ ]: MORTGAGE      197110
RENT          158770
OWN           37443
OTHER         142
Name: home_ownership, dtype: int64
```

```
In [ ]: df['title'].value_counts()[:20]
```

```
Out[ ]: Debt consolidation      152392
Credit card refinancing      51476
Home improvement              15245
Other                        12910
Debt Consolidation           11584
Major purchase                4759
Consolidation                 3840
debt consolidation            3543
Business                     2947
Debt Consolidation Loan       2859
Medical expenses              2733
Car financing                 2135
Credit Card Consolidation     1768
Vacation                      1715
Moving and relocation          1688
consolidation                 1594
Personal Loan                 1568
Consolidation Loan            1295
Home Improvement              1265
Home buying                   1183
Name: title, dtype: int64
```

- (Consolidation, consolidation), (Debt consolidation, debt consolidation) etc are repeating.
- using str.lower to remove such discrepancies

```
In [ ]: df['title'] = df['title'].str.lower()
df['title'].value_counts()[:20]
```

```
Out[ ]: debt consolidation      168000
credit card refinancing      51770
home improvement              17093
other                        12973
consolidation                 5570
major purchase                4988
debt consolidation loan       3508
business                     3015
medical expenses              2811
credit card consolidation     2629
personal loan                 2424
car financing                 2156
credit card payoff            1904
consolidation loan            1881
vacation                      1863
credit card refinance         1832
moving and relocation          1692
consolidate                   1523
```

```
personal      1457
home buying   1196
Name: title, dtype: int64
```


## bar plots

In [ ]:

```
plt.figure(figsize=(20, 8))

plt.subplot(1, 2, 1)
sns.countplot(x = df['grade'], hue=df['loan_status'], order = sorted(df['grade'].unique()))

plt.subplot(1, 2, 2)
sns.countplot(x = df['sub_grade'], hue=df['loan_status'], order = sorted(df['sub_grade'].unique()))
plt.show()
```



Insights:-

1. Grade A, B, C and D are grades where fully-paid applicants are twice of defaulters.
2. Out of above grades, grade B applicants are maximum that paid the loan fully.

In [ ]:

```
plt.figure(figsize=(20, 15))

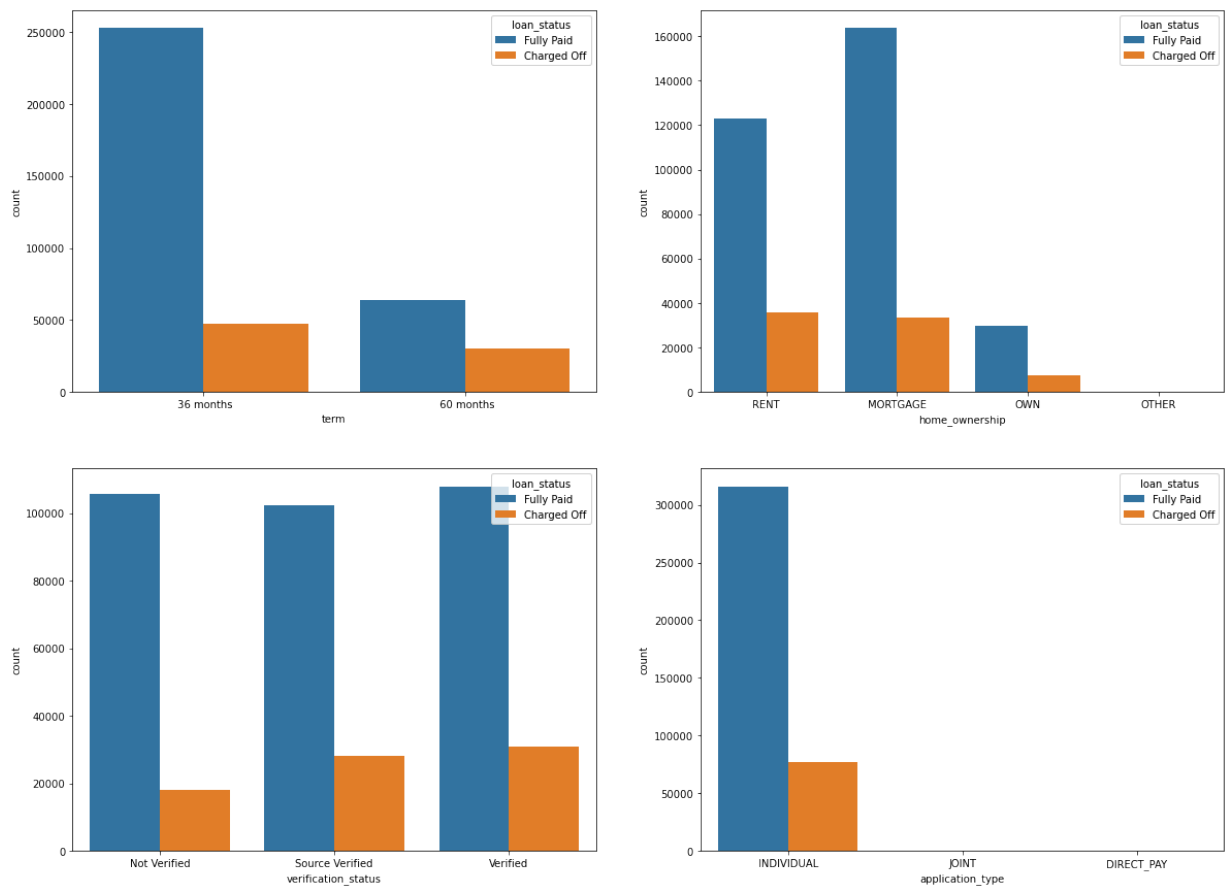
plt.subplot(2, 2, 1)
sns.countplot(x=df['term'], hue=df['loan_status'])

plt.subplot(2, 2, 2)
sns.countplot(x=df['home_ownership'], hue=df['loan_status'])

plt.subplot(2, 2, 3)
sns.countplot(x=df['verification_status'], hue=df['loan_status'])

plt.subplot(2, 2, 4)
sns.countplot(x=df['application_type'], hue=df['loan_status'])

plt.show()
```



### Insights:-

1. 36 months period is opted by most of the applicants.
2. 36 months tenure loan has high success rate (fully paid rate) than 60 months rate.
3. Most of the applicants report their home-ownership as 'Mortgage' or 'Rent'.
4. Individual Applicants generally applied for the loan.

In [ ]:

```
plt.figure(figsize=(15, 10))

plt.subplot(2, 2, 1)
sns.countplot(x=df['pub_rec'], hue=df['loan_status'])

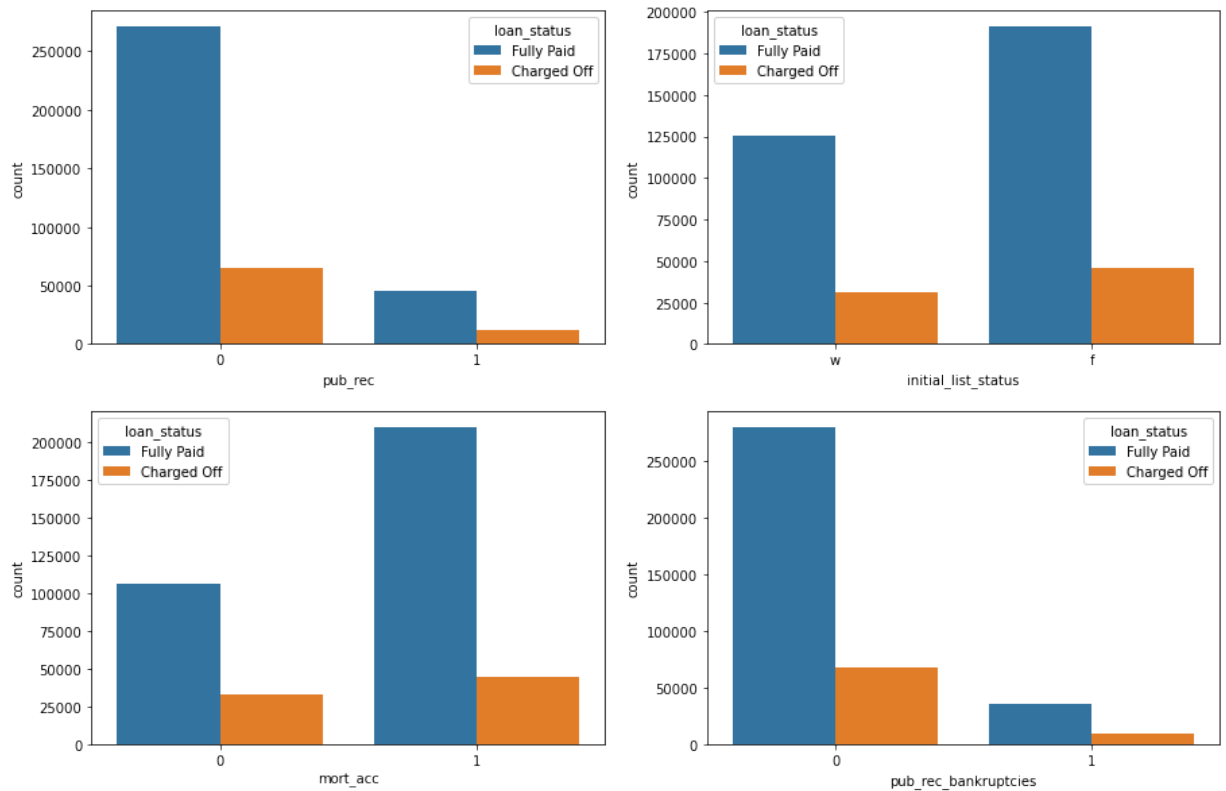
plt.subplot(2, 2, 2)
sns.countplot(x=df['initial_list_status'], hue=df['loan_status'])

plt.subplot(2, 2, 3)
sns.countplot(x=df['mort_acc'], hue=df['loan_status'])

plt.subplot(2, 2, 4)
sns.countplot(x=df['pub_rec_bankruptcies'], hue=df['loan_status'])

plt.show()
```





### Insights:-

1. Applicants with no derogatory remarks have high chance of getting loan and higher success of paying the loan back.
2. fractional loan status count is slightly high compared to whole loan status.
3. Applicants with mortgage accounts have high chance of getting loan and higher success of paying the loan back.
4. Applicants with no bankruptcies have high chance of getting loan and higher success of paying the loan back.


In [ ]:

```
plt.figure(figsize = (20, 8))

plt.subplot(1, 2, 1)
sns.countplot(y = df['purpose'], hue = df['loan_status'])

plt.subplot(1, 2, 2)
order = ['< 1 year', '1 year', '2 years', '3 years', '4 years', '5 years', '6 years']
sns.countplot(y = df['emp_length'], hue = df['loan_status'], order = order)

plt.show()
```



Insights:-

1. Debt consolidation and credit card are the major purpose to be reported in loan application.
2. Applicants with larger length of employment has high chance of getting a loan.


In [ ]:

```
plt.figure(figsize = (20, 8))

plt.subplot(1, 2, 1)
df['emp_title'].value_counts()[1:21].plot(kind='bar')

plt.subplot(1, 2, 2)
df['title'].value_counts()[ :20].plot(kind='bar')
```

Out [ ]: <AxesSubplot:>



Insights:-

1. Teacher and Manager are occupation with highest chance of getting the loan.

In [ ]:


```
plt.figure(figsize = (20, 8))

plt.subplot(1, 2, 1)
sns.countplot(x = df['zip_code'], hue = df['loan_status'])
plt.xticks(rotation = 90)

plt.subplot(1, 2, 2)
top_20_states = df['state'].value_counts()[1:21].index
```

```
df_top_20 = df.loc[df['state'].isin(top_20_states)]
sns.countplot(x = df_top_20['state'], hue=df_top_20['loan_status'], order=top_20_sta
plt.xticks(rotation = 90)

plt.show()
```



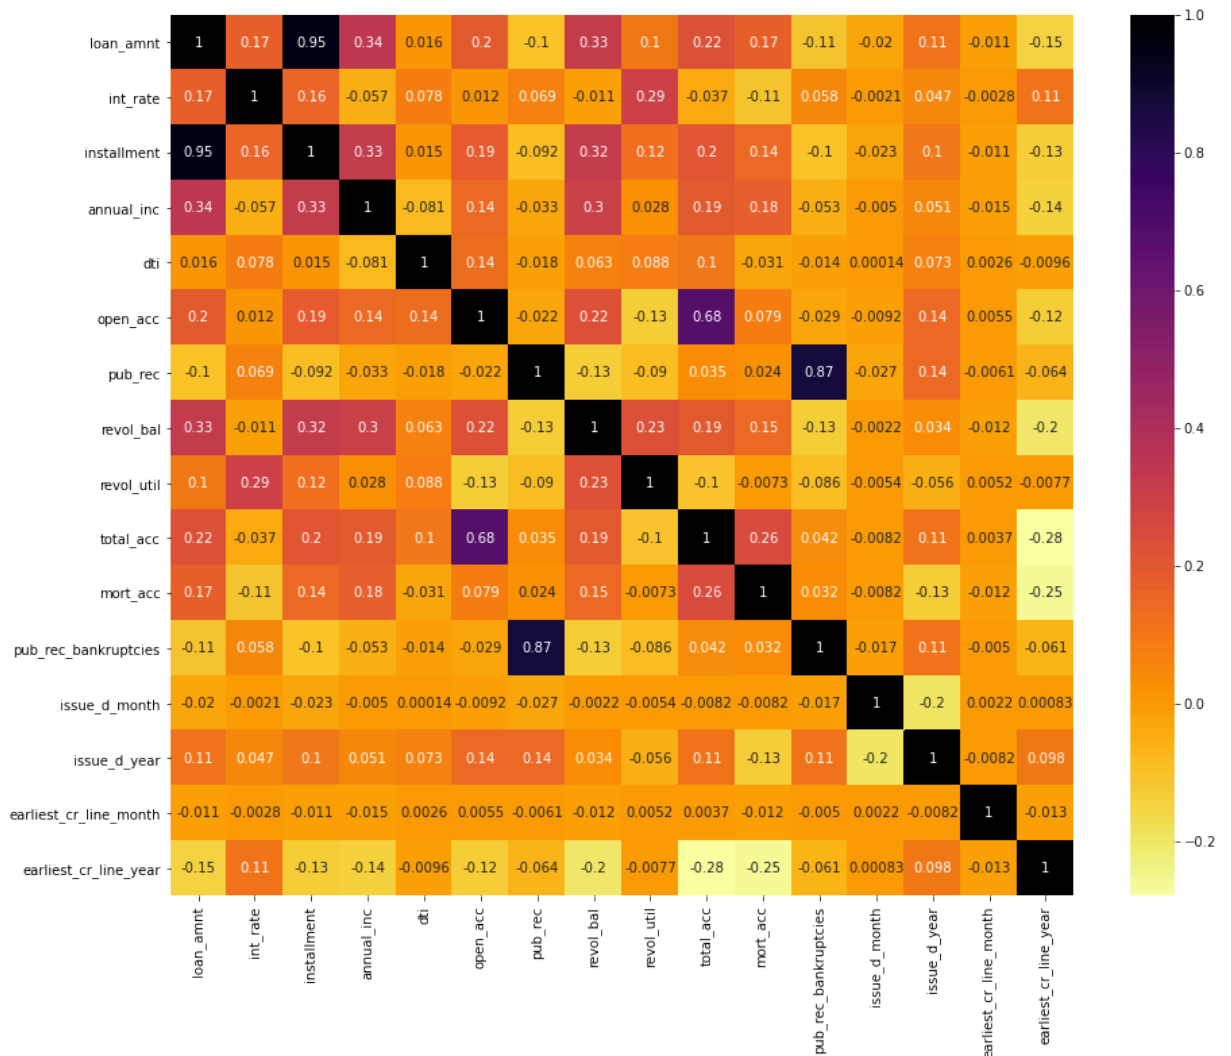
Insights:-

1. Zip-code: 05113, 00813, 29597 are the areas where almost every applicant has paid their loan successfully.
2. Zip-code 11650, 86630, 93700 are the areas where almost every applicant has defaulted.

## Co-relation Matrix and Heatmaps


In [ ]:

```
plt.figure(figsize=(15, 12))
sns.heatmap(df.corr(method="pearson"), annot = True, cmap = "inferno_r")
plt.show()
```



In [ ]:

```
plt.figure(figsize=(15, 12))
sns.heatmap(df.corr(method="spearman"), annot = True, cmap = "inferno_r")
plt.show()
```



- (loan\_amnt, installment), (pub\_rec\_bankruptcies, pub\_rec) pairs are highly co-related.
- Hence dropping installment and pub\_rec\_bankruptcies.

```
In [ ]: df.drop(columns=['installment', 'pub_rec_bankruptcies'], inplace=True)
```


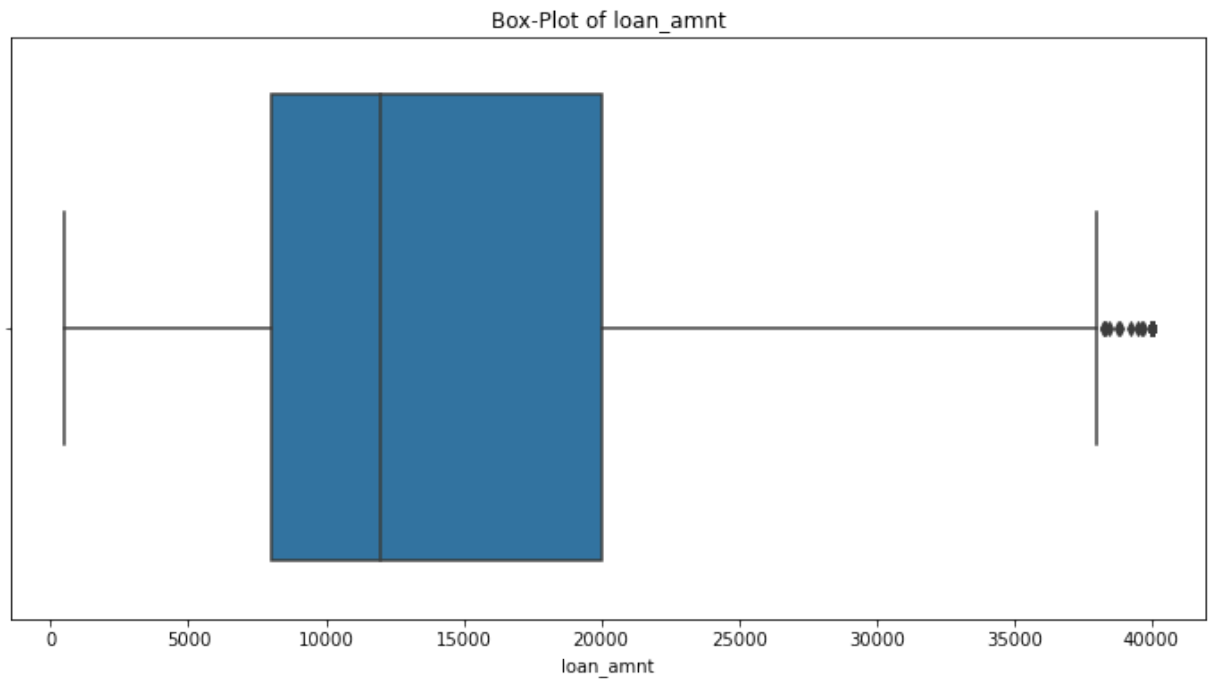
## Outliers Detection using histogram and their treatment


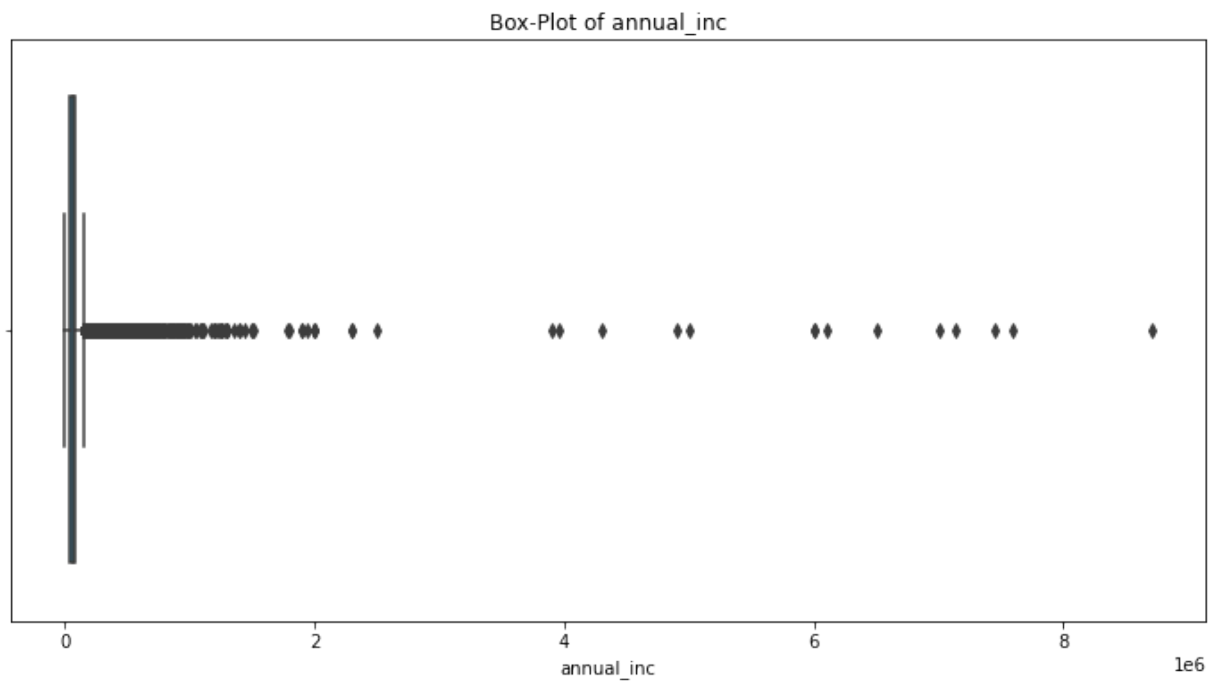
```
In [ ]: num_cols = df.select_dtypes(include='number').columns
num_cols
```

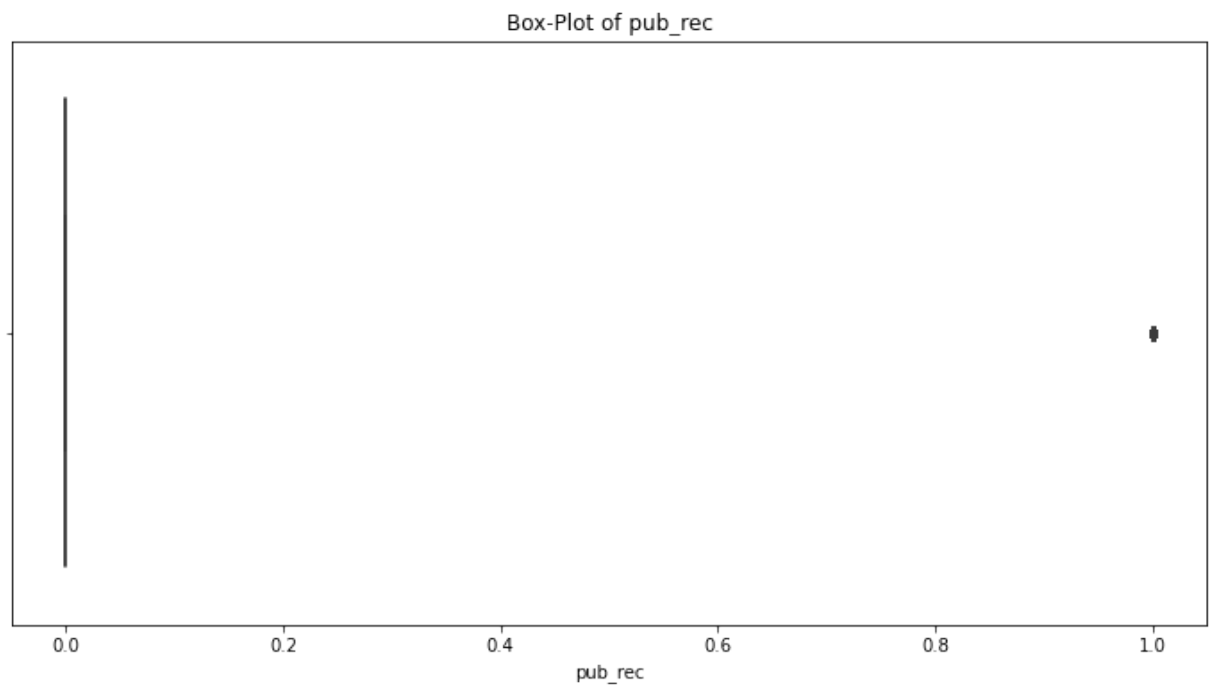
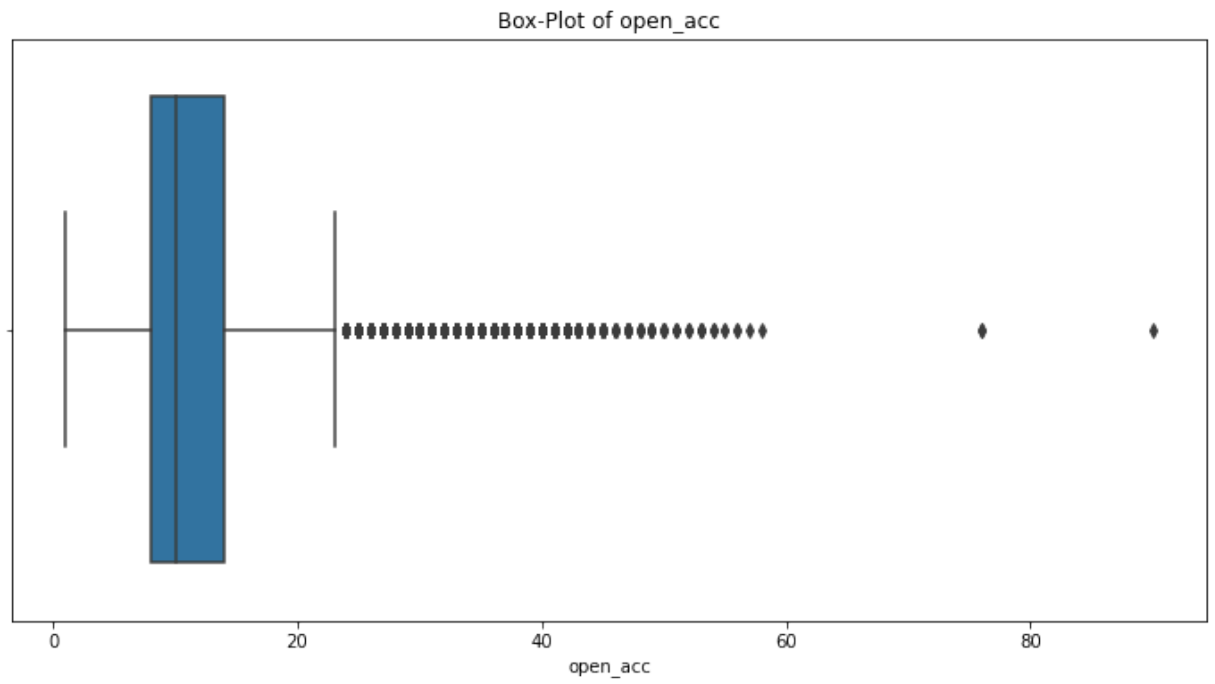
```
Out[ ]: Index(['loan_amnt', 'int_rate', 'annual_inc', 'dti', 'open_acc', 'pub_rec',
              'revol_bal', 'revol_util', 'total_acc', 'mort_acc', 'issue_d_month',
              'issue_d_year', 'earliest_cr_line_month', 'earliest_cr_line_year'],
              dtype='object')
```

```
In [ ]: def show_box_plot(feature: str) -> None:
plt.figure(figsize=(12, 6))
sns.boxplot(x = df[feature])
plt.title(f'Box-Plot of {feature}')
plt.show()
```

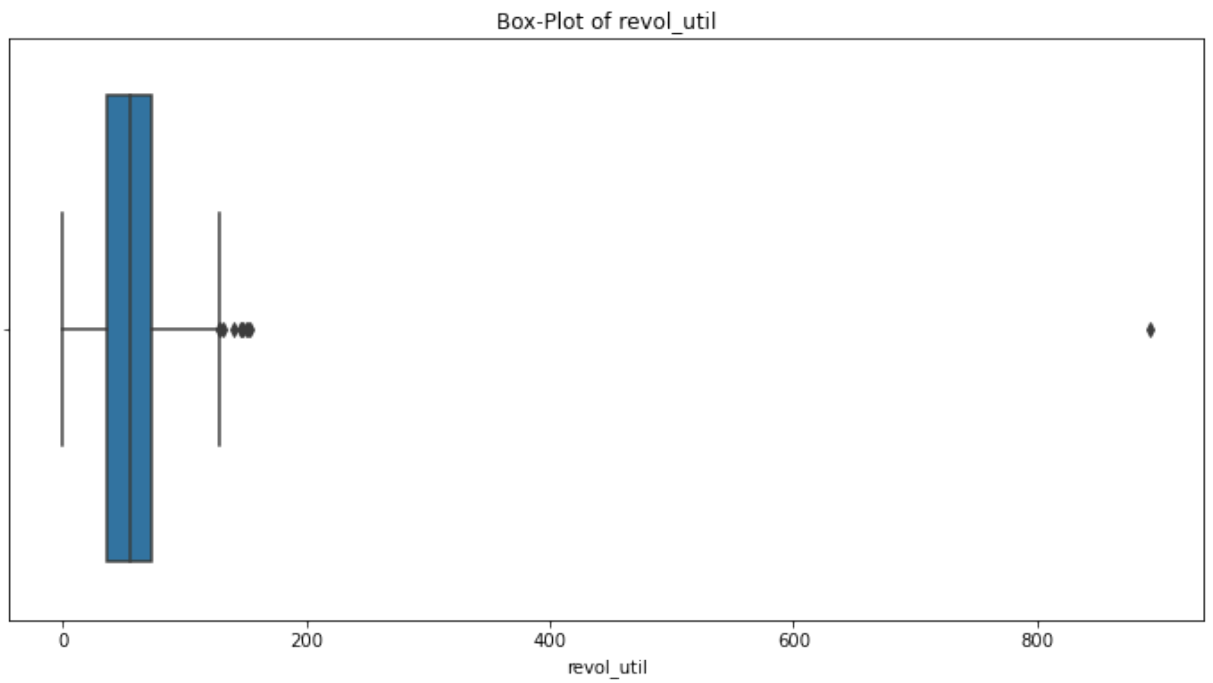
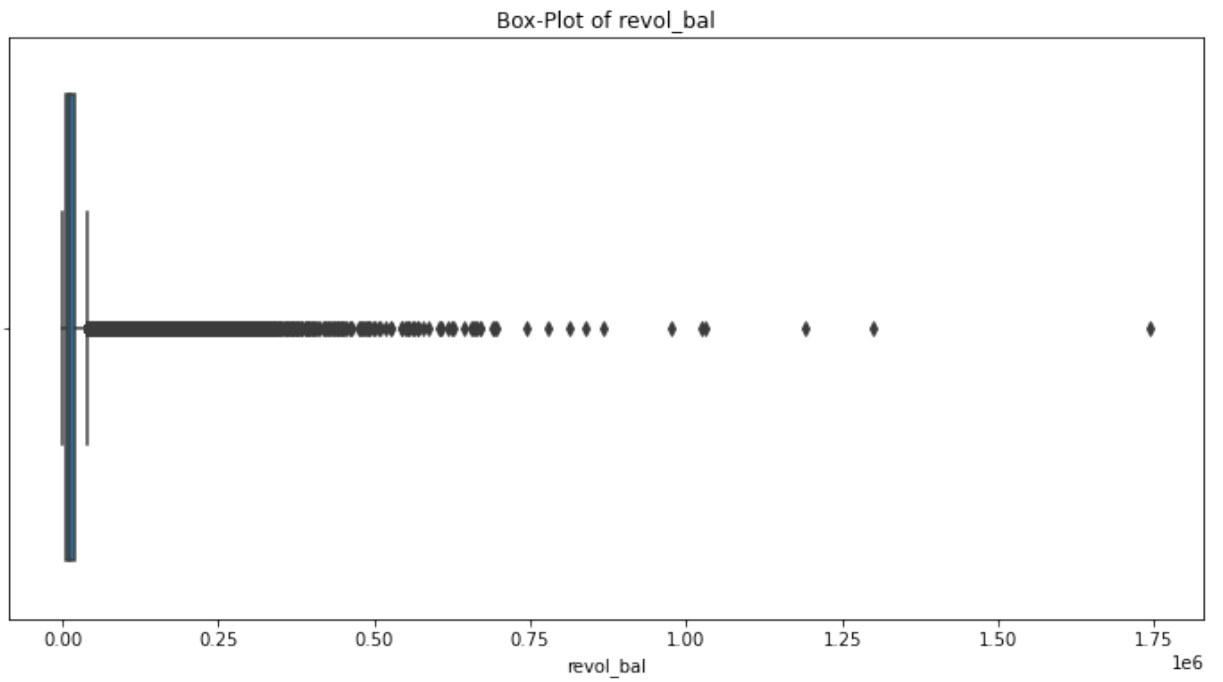
```
In [ ]: for feature in num_cols:
show_box_plot(feature)
```

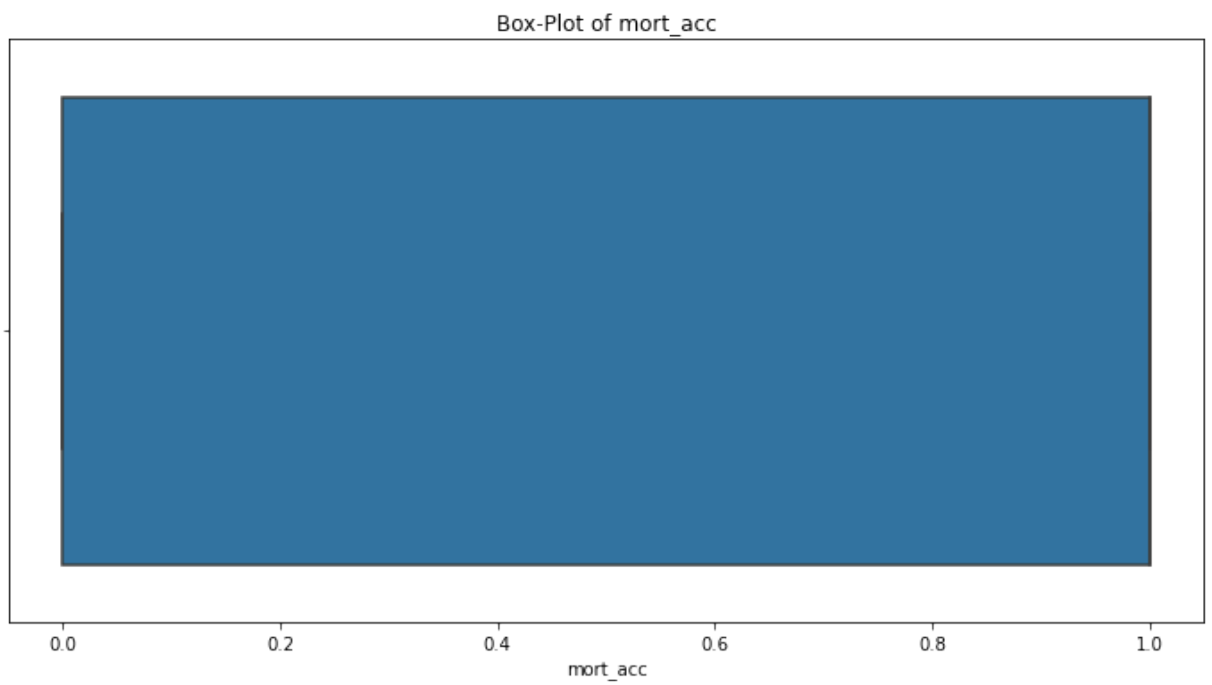



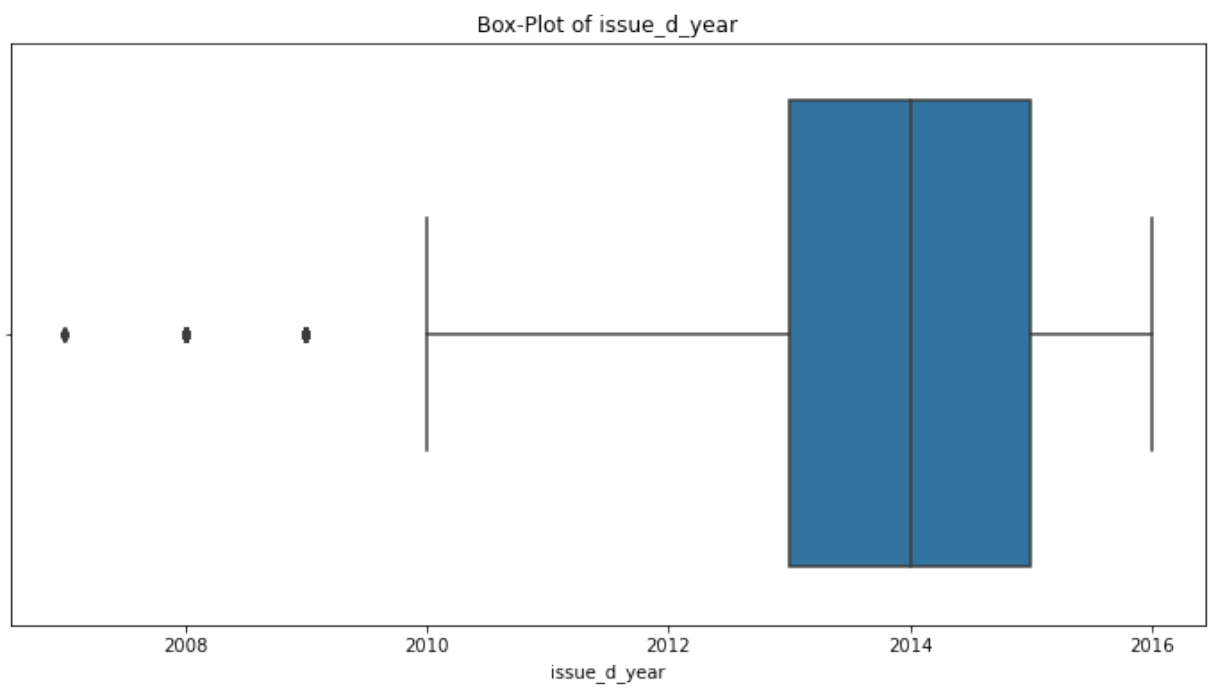
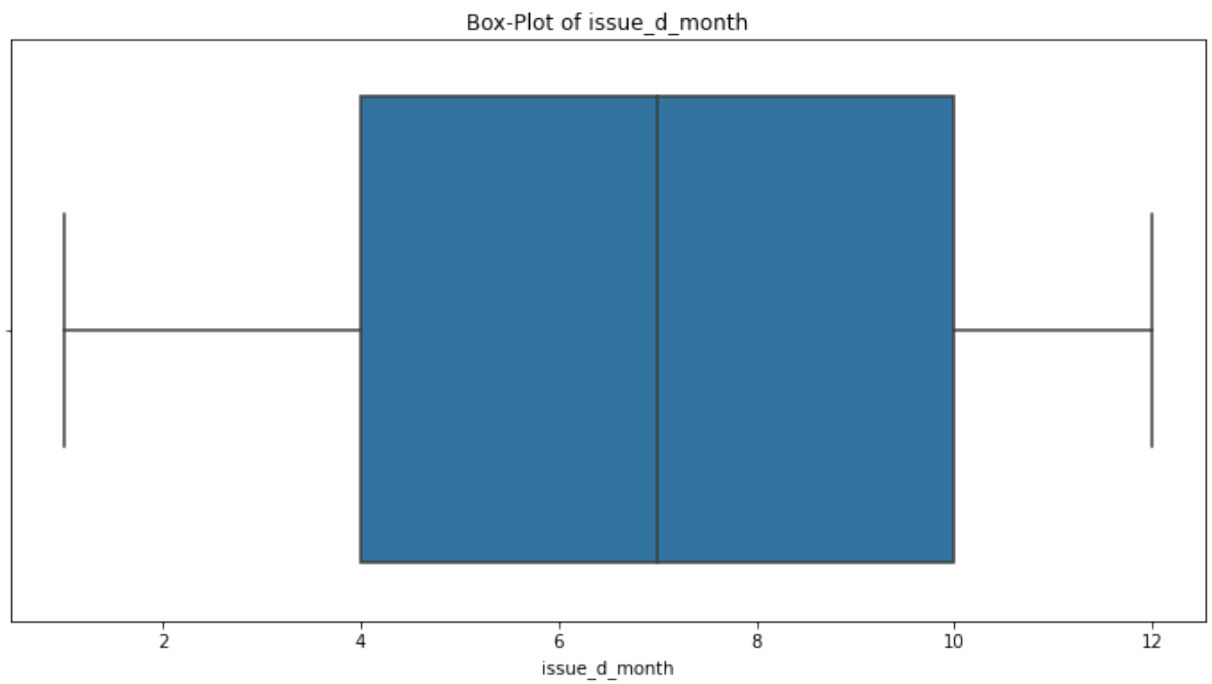















- Boxplots have shown lot of outliers,
- let's treat them using z-score trimming method.

```
In [ ]: for feature in num_cols:
        mean = df[feature].mean()
        std_dev = df[feature].std()

        lower_limit = mean - 3*std_dev
        upper_limit = mean + 3*std_dev

        df = df[ (df[feature] > lower_limit) & (df[feature] < upper_limit) ]
```

```
In [ ]: df.shape
```

```
Out[ ]: (367495, 28)
```

# Encoding

## Label Encoding

```
In [ ]: # 'term' column cleaning
term_mapping = {
    ' 36 months': 36,
    ' 60 months': 60
}

df['term'] = df['term'].map(term_mapping)
df['term'].unique()
```

```
Out[ ]: array([36, 60], dtype=int64)
```

```
In [ ]: # 'grade' column encoding
grade_mapping = {'A': 1, 'B': 2, 'C': 3, 'D': 4, 'E': 5, 'F': 6, 'G': 7}

df['grade'] = df['grade'].map(grade_mapping)
df['grade'].unique()
```

```
Out[ ]: array([2, 1, 3, 5, 4, 6, 7], dtype=int64)
```

```
In [ ]: # 'sub_grade' column encoding
# encoding A1, B1, C1, ... -> 1 (because grade feature will take care of grades)
sub_grade_mapping = {ele: int(ele[1]) for ele in df['sub_grade'].values}

df['sub_grade'] = df['sub_grade'].map(sub_grade_mapping)
df['sub_grade'].unique()
```

```
Out[ ]: array([4, 5, 3, 2, 1], dtype=int64)
```

```
In [ ]: # 'emp_length' column cleaning
emp_length_mapping = {
    '10+ years': 10, '4 years': 4, '< 1 year': 0, '6 years': 6, '9 years': 9,
    '2 years': 2, '3 years': 3, '8 years': 8, '7 years': 7, '5 years': 5, '1 year':
}

df['emp_length'] = df['emp_length'].map(emp_length_mapping)
df['emp_length'].unique()
```

```
Out[ ]: array([10, 4, 0, 6, 9, 2, 3, 7, 8, 5, 1], dtype=int64)
```

### Mapping of Target Feature

```
In [ ]: # loan_status column encoding
# loan_status is the target variable
loan_status_mapping = {'Fully Paid': 0, 'Charged Off': 1}

df['loan_status'] = df['loan_status'].map(loan_status_mapping)
df['loan_status'].unique()
```

```
Out[ ]: array([0, 1], dtype=int64)
```

### Mapping of initial\_list\_status feature

- It contains 2 unique values so Label Encoding is a good choice.

```
In [ ]: df['initial_list_status'].unique()
```

```
Out[ ]: array(['w', 'f'], dtype=object)
```

```
In [ ]: df['initial_list_status'] = df['initial_list_status'].map({'w': 0, 'f': 1})
df['initial_list_status'].unique()
```

```
Out[ ]: array([0, 1], dtype=int64)
```

## One-hot Encoding

```
In [ ]: df.select_dtypes(include="object").nunique()
```

```
Out[ ]: emp_title      162992
home_ownership      4
verification_status  3
purpose             14
title               37131
application_type     3
state               52
zip_code            10
dtype: int64
```

- emp\_title, title, state contains lots of unique values, so I am dropping them to avoid curse of dimensionality.

```
In [ ]: df.drop(['emp_title', 'title', 'state'], axis=1, inplace=True)
```

Performing one-hot encoding now for other features


```
In [ ]: dummies = ['purpose', 'zip_code', 'verification_status', 'application_type', 'home_o
df = pd.get_dummies(df, columns=dummies, drop_first=True)
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

|   | loan_amnt | term | int_rate | grade | sub_grade | emp_length | annual_inc | loan_status | dti   | open_ac |
|---|-----------|------|----------|-------|-----------|------------|------------|-------------|-------|---------|
| 0 | 10000.0   | 36   | 11.44    | 2     | 4         | 10         | 117000.0   | 0           | 26.24 | 16.     |
| 1 | 8000.0    | 36   | 11.99    | 2     | 5         | 4          | 65000.0    | 0           | 22.05 | 17.     |
| 2 | 15600.0   | 36   | 10.49    | 2     | 3         | 0          | 43057.0    | 0           | 12.79 | 13.     |
| 3 | 7200.0    | 36   | 6.49     | 1     | 2         | 6          | 54000.0    | 0           | 2.60  | 6.      |
| 4 | 24375.0   | 60   | 17.27    | 3     | 5         | 9          | 55000.0    | 1           | 33.95 | 13.     |

5 rows × 49 columns



## Saving the data to a file


```
In [ ]: df.to_csv("loan_tap_clean.csv", index=False)
```

## Building Classification Model

```
In [ ]: df = pd.read_csv('loan_tap_clean.csv')
df.head()
```

```
Out[ ]:      loan_amnt  term  int_rate  grade  sub_grade  emp_length  annual_inc  loan_status  dti  open_ac
0      10000.0   36    11.44     2        4           10    117000.0           0  26.24     16.
1       8000.0   36    11.99     2        5           4     65000.0           0  22.05     17.
2      15600.0   36    10.49     2        3           0     43057.0           0  12.79     13.
3       7200.0   36     6.49     1        2           6     54000.0           0   2.60      6.
4      24375.0   60    17.27     3        5           9     55000.0           1  33.95     13.
```

5 rows × 49 columns



## Utility functions

```
In [ ]: def show_roc_curve(y_true, y_pred_proba):
    fpr, tpr, thresholds = metrics.roc_curve(y_true, y_pred_proba)
    roc_area_score = metrics.auc(fpr, tpr)

    plt.figure(figsize=(8, 6))

    plt.plot(fpr, tpr, label = f'Logistic Regression (ROC Area Score: {round(roc_area_score, 2)})')
    plt.plot(fpr, fpr, 'r--', label = "Random Model (ROC Area Score: 0.5)")

    # Axis Labels
    plt.xlabel("False Positive rate")
    plt.ylabel("True Positive Rate")

    # Other properties
    plt.grid()
    plt.title("ROC curve")
    plt.legend(loc = "lower right")

    plt.show()
```

```
In [ ]: def show_pr_curve(y_true, y_pred_proba):
    precisions, recalls, thresholds = metrics.precision_recall_curve(y_true, y_pred_proba)
    pr_area_score = metrics.auc(recalls, precisions)

    plt.figure(figsize=(8, 6))

    plt.plot(recalls, precisions, label = f'Logistic Regression (PR score: {round(pr_area_score, 2)})')
    plt.plot(recalls, 1 - recalls, 'r--', label = "Random Model (ROC score: 0.5)")

    # Axis Labels
    plt.xlabel("Recalls")
```

```
plt.ylabel("Precisions")

# Other properties
plt.title("PR curve")
plt.legend(loc = "lower left")
plt.grid()

plt.show()
```

In [ ]:

```
def show_pr_trade_off_curve(y_true, y_pred_proba):
    precisions, recalls, thresholds = metrics.precision_recall_curve(y_true, y_pred_
    boundary = thresholds.shape[0]

    plt.figure(figsize=(8, 6))

    # precision plot
    plt.plot(thresholds, precisions[:boundary], 'b--', label = 'precision')

    # recall plot
    plt.plot(thresholds, recalls[:boundary], 'r--', label = 'recall')

    # Axis Labels
    plt.xlabel("Threshold")
    plt.xlabel("Precision/Recall")

    # Other properties
    plt.legend()
    plt.grid()
    plt.title('Precision-Recall Trade off')

    plt.show()
```

In [ ]:

```
def show_metrics_summary(y_true, y_pred, y_pred_proba):
    print(metrics.classification_report(y_true, y_pred))
    print('-----')

    cm = metrics.confusion_matrix(y_true, y_pred)
    print(cm)
    print('-----')
    print(cm*100/cm.sum())
    print('-----')

    metrics.ConfusionMatrixDisplay(cm).plot()
    plt.show()

    show_roc_curve(y_true, y_pred_proba)
    show_pr_curve(y_true, y_pred_proba)
    show_pr_trade_off_curve(y_true, y_pred_proba)
```

## Logistic Regression Model -1

- without class balancing, regularizer and multi-collinearity check

### X and y

In [ ]:

```
X = df.drop('loan_status', axis=1)
y = df['loan_status']
```



**train-test split**

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10, test_size=0.3)

print(X_train.shape)
print(X_test.shape)
```

(220497, 48)  
(146998, 48)

**Scaling**

```
In [ ]: scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

**Classification Model**

```
In [ ]: model = LogisticRegression(max_iter=1000) # Adjust iterations based on need

model.fit(X_train, y_train)
```

```
Out[ ]: LogisticRegression
LogisticRegression(max_iter=1000)
```


**Model Performace**


```
In [ ]: y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1]
```

```
In [ ]: show_metrics_summary(y_test, y_pred, y_pred_proba)
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.89      | 0.99   | 0.93     | 117931  |
| 1            | 0.91      | 0.48   | 0.63     | 29067   |
| accuracy     |           |        | 0.89     | 146998  |
| macro avg    | 0.90      | 0.73   | 0.78     | 146998  |
| weighted avg | 0.89      | 0.89   | 0.87     | 146998  |

```
-----
--
[[116581  1350]
 [ 15082 13985]]
-----
--
[[79.30788174  0.91837984]
 [10.26000354  9.51373488]]
-----
--
```





```
In [ ]: metrics.fbeta_score(y_test, y_pred, beta=2)
```

```
Out[ ]: 0.5313328723509343
```

Insights:-

1. Metrics summary of Logistic Regression without handling balancing, regularization and multi-collinearity check:

- Accuracy: 0.89
- Precision: 0.91
- Recall: 0.48
- F1-score: 0.63
- ROC-AUC: 0.91
- PR-AUC: 0.78
- fbeta(beta=2): 0.53

2. Confusion Matrix:

- TN: 79.3%
- FP: 0.91%
- FN: 10.26%
- TP: 9.51%

## Multi-Collinearity check and its removal

```
In [ ]: X = df.drop('loan_status', axis=1)
        y = df['loan_status']
```

```
In [ ]: def remove_features_by_vif(X: pd.DataFrame, vif_threshold = 10):
        features = X.columns.to_list()
```

```

dropped_features = list()

while True:
    # calculate vifs
    vifs = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
    vifs = list(zip(features, vifs))

    # sort vifs
    vifs.sort(key= lambda x: x[1], reverse=True)

    # top vif
    top_vif = vifs[0][1]
    top_vif_feature = vifs[0][0]

    if top_vif < vif_threshold:
        break

    # drop top vif feature
    dropped_features.append(top_vif_feature)
    X.drop(columns=[top_vif_feature], inplace=True)
    features.remove(top_vif_feature)

return dropped_features

```

```

In [ ]: dropped_features = remove_features_by_vif(X, vif_threshold= 10)
        dropped_features

```

```

Out[ ]: ['issue_d_year',
        'earliest_cr_line_year',
        'int_rate',
        'application_type_INDIVIDUAL',
        'term',
        'purpose_debt_consolidation',
        'open_acc']

```

## Logistic Regression after removal of multi-collinearity

### train-test split

```

In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10, test_size

print(X_train.shape)
print(X_test.shape)

```

```

(220497, 41)
(146998, 41)

```

### Scaling

```

In [ ]: scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

### Classification Model

```

In [ ]: model = LogisticRegression(max_iter=1000) # Adjust iterations based on need

model.fit(X_train, y_train)

```

Out [ ]: **LogisticRegression**  
 LogisticRegression(max\_iter=1000)

### Model Performace


In [ ]: `y_pred = model.predict(X_test)`  
`y_pred_proba = model.predict_proba(X_test)[:, 1]`


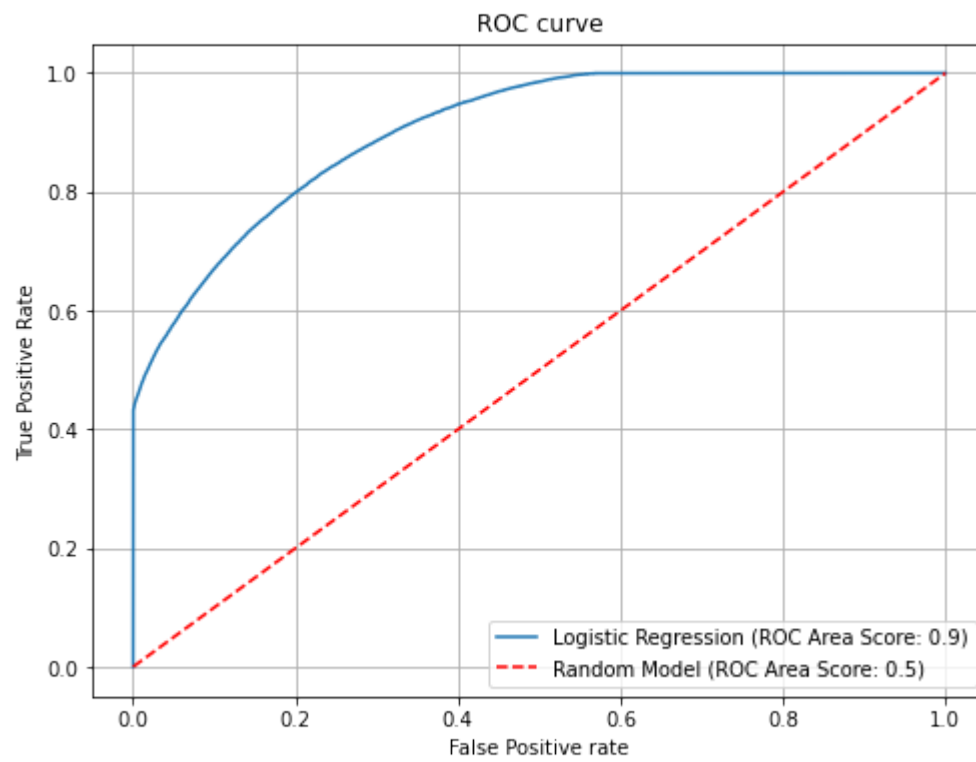
In [ ]: `show_metrics_summary(y_test, y_pred, y_pred_proba)`


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.88      | 0.99   | 0.93     | 117931  |
| 1            | 0.93      | 0.47   | 0.63     | 29067   |
| accuracy     |           |        | 0.89     | 146998  |
| macro avg    | 0.90      | 0.73   | 0.78     | 146998  |
| weighted avg | 0.89      | 0.89   | 0.87     | 146998  |

```
--
[[116820  1111]
 [ 15335 13732]]
--
```

```
--
[[79.47046899  0.7557926 ]
 [10.43211472  9.3416237 ]]
--
```







```
In [ ]: metrics.fbeta_score(y_test, y_pred, beta=2)
```

```
Out[ ]: 0.5236784098969576
```

Insights:-

1. Metrics summary of Logistic Regression after multi-collinearity check:

- Accuracy: 0.89
- Precision: 0.93
- Recall: 0.47
- F1-score: 0.63
- ROC-AUC: 0.9
- PR-AUC: 0.78
- fbeta(beta=2): 0.52

2. Confusion Matrix:

- TN: 79.47%
- FP: 0.75%
- FN: 10.43%
- TP: 9.34%

## Handling Imbalance using class-weights

**X and y**

```
In [ ]: X = df.drop('loan_status', axis=1)
        y = df['loan_status']
```

**Drop non-important features**

```
In [ ]: X.drop(columns=dropped_features, inplace=True)
```

### train-test split

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10, test_size

print(X_train.shape)
print(X_test.shape)
```

```
(220497, 41)
```

```
(146998, 41)
```

### Scaling

```
In [ ]: scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

### Classification Model

```
In [ ]: model = LogisticRegression(max_iter=1000, class_weight='balanced') # Adjust iterati

model.fit(X_train, y_train)
```

```
Out[ ]: LogisticRegression
LogisticRegression(class_weight='balanced', max_iter=1000)
```

### Model Performace

```
In [ ]: y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1]
```

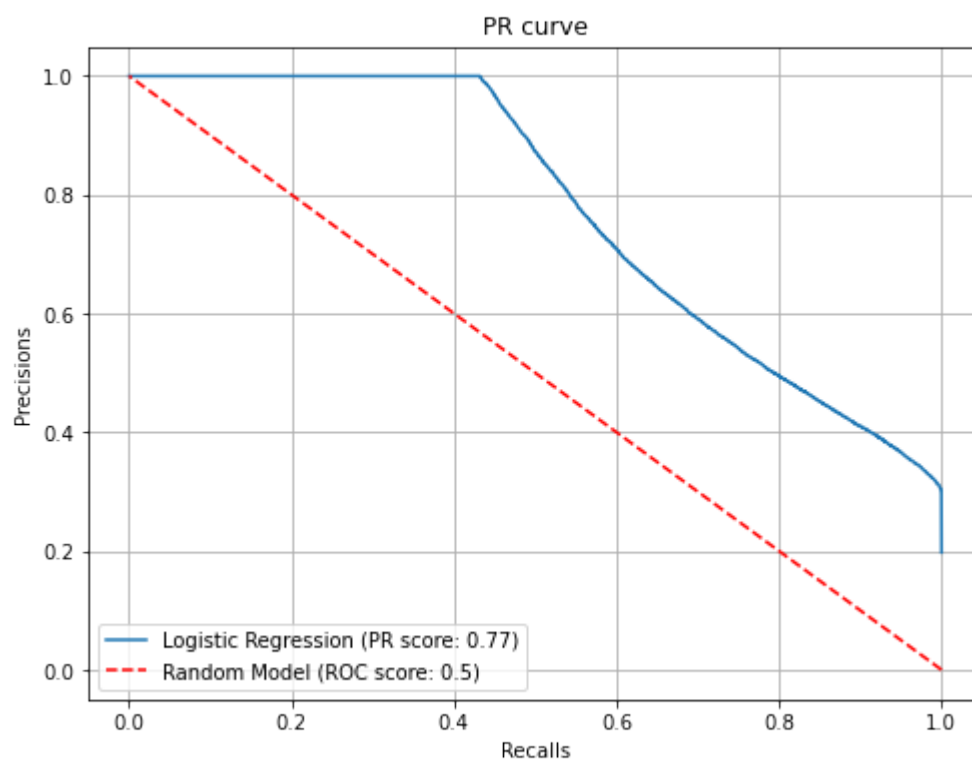


```
In [ ]: show_metrics_summary(y_test, y_pred, y_pred_proba)
```


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.94      | 0.81   | 0.87     | 117931  |
| 1            | 0.50      | 0.79   | 0.61     | 29067   |
| accuracy     |           |        | 0.80     | 146998  |
| macro avg    | 0.72      | 0.80   | 0.74     | 146998  |
| weighted avg | 0.85      | 0.80   | 0.82     | 146998  |

```
-----
--
[[94950 22981]
 [ 6021 23046]]
-----
```

```
--
[[64.59271555 15.63354603]
 [ 4.09597409 15.67776432]]
-----
--
```







```
In [ ]: metrics.fbeta_score(y_test, y_pred, beta=2)
```

```
Out[ ]: 0.7100033888906004
```

Insights:-

1. Metrics summary of Logistic Regression after handling balancing and multi-collinearity check:

- Accuracy: 0.8
- Precision: 0.5
- Recall: 0.79
- F1-score: 0.61
- ROC-AUC: 0.9
- PR-AUC: 0.77
- fbeta(beta=2): 0.71

2. Confusion Matrix:

- TN: 64.59%
- FP: 15.63%
- FN: 4.09%
- TP: 15.67%

## Regularization (L2 regularizer) and Cross Validation

**X and y**

```
In [ ]: X = df.drop('loan_status', axis=1)
        y = df['loan_status']
```

**Drop non-important features**

```
In [ ]: X.drop(columns=dropped_features, inplace=True)
```

### train-val-test split

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10, test_size=0.2)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, random_state=10, test_size=0.2)

print(X_train.shape)
print(X_val.shape)
print(X_test.shape)
```

```
(183747, 41)
```

```
(91874, 41)
```

```
(91874, 41)
```

### scaling

```
In [ ]: scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)
```

### cross-validation

```
In [ ]: lamdas = [0.01, 0.1, 1, 10]
train_scores = list()
val_scores = list()

for lamda in lamdas:
    model = LogisticRegression(max_iter=1000, class_weight='balanced', penalty='l2',

    # train the model
    model.fit(X_train, y_train)

    # training score
    y_train_pred = model.predict(X_train)
    train_scores.append(metrics.recall_score(y_train, y_train_pred))

    # validation score
    y_val_pred = model.predict(X_val)
    val_scores.append(metrics.recall_score(y_val, y_val_pred))

print(lamdas)
print(train_scores)
print(val_scores)
```

```
[0.01, 0.1, 1, 10]
```

```
[0.7946626719928983, 0.7946071904127829, 0.794329782512206, 0.7925821127385708]
```

```
[0.7881546476556074, 0.7881546476556074, 0.7878256100904854, 0.7863997806416233]
```

- Not much difference in using different lambdas
- let use lambda = 0.01

## Test score/summary

### Classification Model

```
In [ ]: model = LogisticRegression(max_iter=1000, class_weight='balanced', penalty='l2', C=
model.fit(X_train, y_train)
```

```
Out[ ]: LogisticRegression
LogisticRegression(C=100.0, class_weight='balanced', max_iter=1000)
```

### Model Performace


```
In [ ]: y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1]
```



```
In [ ]: show_metrics_summary(y_test, y_pred, y_pred_proba)
```


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.94      | 0.81   | 0.87     | 73683   |
| 1            | 0.50      | 0.79   | 0.62     | 18191   |
| accuracy     |           |        | 0.80     | 91874   |
| macro avg    | 0.72      | 0.80   | 0.74     | 91874   |
| weighted avg | 0.85      | 0.80   | 0.82     | 91874   |

```
--
[[59553 14130]
 [ 3833 14358]]
--
```

```
--
[[64.82029736 15.37975924]
 [ 4.1720182  15.6279252  ]]
--
```







```
In [ ]: metrics.fbeta_score(y_test, y_pred, beta=2)
```

```
Out[ ]: 0.7090230316438193
```

Insights:-

1. Metrics summary of Logistic Regression after handling balancing, multi-collinearity check, cross-validation and regularization:

- Accuracy: 0.8
- Precision: 0.5
- Recall: 0.79
- F1-score: 0.62
- ROC-AUC: 0.9
- PR-AUC: 0.77
- fbeta(beta=2): 0.71

2. Confusion Matrix:

- TN: 64.82%
- FP: 15.37%
- FN: 4.17%
- TP: 15.62%

## Logistic Regression Model - (final)

**X and y**

```
In [ ]: X = df.drop('loan_status', axis=1)
        y = df['loan_status']
```

**Drop non-important features**

```
In [ ]: X.drop(columns=dropped_features, inplace=True)
```

### train-test split

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10, test_size

print(X_train.shape)
print(X_test.shape)
```

```
(220497, 41)
```

```
(146998, 41)
```

### Scaling

```
In [ ]: scaler = MinMaxScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

### Classification Model

```
In [ ]: model = LogisticRegression(max_iter=1000, class_weight='balanced', penalty='l2', C=1

model.fit(X_train, y_train)
```

```
Out[ ]: LogisticRegression
LogisticRegression(C=100.0, class_weight='balanced', max_iter=1000)
```

### Model Performace


```
In [ ]: y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1]
```

```
In [ ]: show_metrics_summary(y_test, y_pred, y_pred_proba)
```


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.94      | 0.80   | 0.87     | 117931  |
| 1            | 0.50      | 0.79   | 0.61     | 29067   |
| accuracy     |           |        | 0.80     | 146998  |
| macro avg    | 0.72      | 0.80   | 0.74     | 146998  |
| weighted avg | 0.85      | 0.80   | 0.82     | 146998  |

```
--
[[94928 23003]
 [ 6019 23048]]
--
```

```
--
[[64.57774936 15.64851222]
 [ 4.09461353 15.67912489]]
--
```







```
In [ ]: metrics.fbeta_score(y_test, y_pred, beta=2)
```

```
Out[ ]: 0.7099600170035547
```

### comparing weights

```
In [ ]: feature_with_weights = list(zip(X.columns, model.coef_.flatten().tolist()))
feature_with_weights.sort(key=lambda x: x[1], reverse=True)

feature_with_weights[:10]
```

```
Out[ ]: [('zip_code_93700', 20.507644059351133),
('zip_code_11650', 20.49416052939432),
('zip_code_86630', 20.35131438231251),
('zip_code_48052', 9.785400878146039),
('zip_code_70466', 9.753018661142576),
('zip_code_30723', 9.747578592989196),
('zip_code_22690', 9.734376695171463),
('grade', 2.6427442289162046),
('dti', 1.7199235354775095),
('loan_amnt', 0.7445231314973701)]
```

Insights:-

1. Metrics summary of my final Logistic Regression model:

- Accuracy: 0.8
- Precision: 0.5
- Recall: 0.79
- F1-score: 0.61
- ROC-AUC: 0.9
- PR-AUC: 0.77
- fbeta(beta=2): 0.71

2. Confusion Matrix:

- TN: 64.57%
- FP: 15.64%
- FN: 4.09%
- TP: 15.67%

3. Zip-code (address), grade, dti and loan\_amnt are the most important features by weights.

# Conclusion

## Summarizing Insights

### Exploratory Data Analysis (EDA) Insights:

- **Dataset Overview:**
  - Total Rows: 396,030
  - Total Columns: 27
  - No duplicates; some missing values were noted.
- **Grade and Repayment:**
  - Grades A, B, C, and D have a higher number of fully-paid applicants compared to defaulters.
  - Grade B has the highest count of fully-paid loans.
- **Loan Term Insights:**
  - The 36-month period is the most common choice among applicants.
  - Loans with a 36-month term have a higher success rate compared to 60-month loans.
- **Applicant Characteristics:**
  - Most applicants report home ownership as 'Mortgage' or 'Rent.'
  - Individual applicants are more likely to apply for loans.
  - Applicants with no derogatory remarks or bankruptcies, and those with mortgage accounts, have a higher chance of receiving loans and successfully repaying them.
- **Loan Purpose:**
  - Debt consolidation and credit card repayment are the most common reasons for applying.
  - Applicants with longer employment histories have a higher likelihood of loan approval.
- **Occupation and Location Insights:**
  - Teachers and Managers are the occupations most likely to be granted loans.
  - Certain ZIP codes (e.g., 05113, 00813) are associated with high loan repayment rates, while others (e.g., 11650, 86630) have high default rates.
- **Feature Correlations:**
  - Strong correlations were found between (loan\_amnt, installment) and (pub\_rec\_bankruptcies, pub\_rec) .

### Logistic Regression Model Insights:

- **Accuracy:** 0.8
- **Precision:** 0.5
- **Recall:** 0.79
- **F1-Score:** 0.61
- **ROC-AUC:** 0.9
- **PR-AUC:** 0.77
- **Confusion Matrix:** TN: 64.57%, FP: 15.64%, FN: 4.09%, TP: 15.67%
- **Key Features:** ZIP code, grade, debt-to-income ratio (DTI), and loan amount were identified as the most important features based on their weights.

## Some Questions

### 1. Comment on Final Model Statistics (including confusion matrix):

The final logistic regression model demonstrates a balanced performance with a strong emphasis on recall. The key statistics are as follows:

- **Accuracy (0.8):** The model correctly predicts 80% of the cases, indicating a good overall performance.
- **Precision (0.5):** Precision is moderate, reflecting the trade-off between false positives and true positives. This indicates that while the model identifies a reasonable number of actual defaulters, it also misclassifies a notable proportion of non-defaulters as defaulters.
- **Recall (0.79):** The recall is high, which is crucial for minimizing false negatives—cases where a defaulter is incorrectly classified as a non-defaulter. This metric is particularly important in credit risk scenarios, where missing a defaulter can lead to significant financial losses.
- **F1-Score (0.61):** The F1-score balances precision and recall, providing a single metric to evaluate the model's effectiveness. The score of 0.61 indicates a reasonable balance between the two.
- **ROC-AUC (0.9):** The high ROC-AUC value suggests that the model has a strong ability to distinguish between defaulters and non-defaulters.
- **PR-AUC (0.77):** The PR-AUC score further confirms the model's effectiveness in handling the imbalanced nature of the dataset.
- **Fbeta (0.71):** With a beta value of 2, the fbeta score prioritizes recall over precision, and a score of 0.71 reflects the model's strength in correctly identifying defaulters.

The confusion matrix reveals a higher true negative rate (64.57%) compared to the true positive rate (15.67%), indicating the model is more conservative in predicting defaults. The false positive rate (15.64%) is slightly higher than desired, which suggests room for improvement in precision. However, the low false negative rate (4.09%) highlights the model's effectiveness in minimizing missed defaults, aligning with the business goal of reducing financial risk.

Overall, this model provides a robust framework for LoanTap to make informed lending decisions, balancing the need to identify high-risk applicants while minimizing potential losses.

### 2. ROC-AUC curve and its comment

#### Comment on ROC Curve

- The ROC curve illustrates the True Positive Rate (TPR) against the False Positive Rate (FPR) for various threshold values.
- The area under the ROC curve (AUC) is 0.9, which signifies that the logistic regression model has a strong ability to differentiate between the positive (defaulter) and negative (non-defaulter) classes.

### 3. PR curve and its comment

#### Comment on PR Curve

- The PR curve is particularly useful for evaluating model performance in cases of class imbalance.
- The area under the PR curve (PR AUC) is 0.77, indicating that the model effectively balances precision and recall.

#### Comment on PR Trade-Off Curve

- The PR trade-off curve illustrates the inverse relationship between precision and recall.
- As precision increases, recall typically decreases, and vice versa.
- The curve demonstrates that at various thresholds, the model can achieve either high precision with a trade-off in recall or high recall with a trade-off in precision.
- Understanding this trade-off is crucial for selecting the optimal threshold based on whether the priority is minimizing false positives (favoring precision) or minimizing false negatives (favoring recall).

### 4. How can we make sure that our model can detect real defaulters and there are less false positives? This is important as we can lose out on an opportunity to finance more individuals and earn interest on it.

#### Ensuring the Model Accurately Detects Real Defaulters and Minimizes False Positives

To improve the detection of real defaulters while reducing false positives, several strategies can be employed:

##### 1. Adjust the Decision Threshold:

- The default threshold for classifying defaulters might be 0.5, but this can be adjusted to balance recall and precision. By lowering the threshold, the model can become more sensitive to potential defaulters, reducing false negatives. Conversely, raising the threshold can help minimize false positives.

##### 2. Incorporate More Granular Features:

- Consider adding or engineering features that provide more detailed insights into an applicant's financial behavior, such as credit utilization ratios, recent spending patterns, or more specific employment stability indicators. These can help the model differentiate between high-risk and low-risk applicants more effectively.

##### 3. Implement a Cost-Sensitive Learning Approach:

- Introduce a cost-sensitive learning framework where the cost of misclassifying a non-defaulter (false positive) is higher. This would encourage the model to be more cautious in predicting defaulters, reducing the number of false positives.

#### 4. Use Ensemble Methods:

- Combine multiple models, such as a random forest or gradient boosting, with logistic regression to leverage the strengths of each. Ensemble methods often provide better generalization and can reduce both false positives and false negatives.

#### 5. Apply More Rigorous Cross-Validation:

- Ensure that the model is tested across various subsets of the data through techniques like k-fold cross-validation. This helps in identifying any biases or inconsistencies in the model's predictions, leading to a more reliable and accurate model.

#### 6. Regularly Update the Model with New Data:

- Continuously update the model with the latest data to ensure it captures recent trends and changes in borrower behavior. A model trained on outdated data might not perform well on new applicants, leading to higher false positives.

#### 7. Monitor and Fine-Tune the Model Post-Deployment:

- After deployment, keep track of the model's performance metrics, especially the false positive rate. Fine-tune the model based on real-world data and outcomes to ensure it continues to meet the desired balance between detecting defaulters and minimizing false positives.

#### 8. Leverage Alternative Data Sources:

- Incorporate alternative data sources such as social media activity, digital footprints, or utility payment history. These can provide additional context that traditional credit metrics might miss, helping to refine the model's predictions.

#### 9. Develop a Two-Stage Model:

- Consider implementing a two-stage model where the first stage identifies potential defaulters with high recall, and the second stage refines these predictions with a focus on minimizing false positives. This approach can help balance the competing objectives of recall and precision.

#### 10. Involve Human Review for Borderline Cases:

- For cases where the model's prediction is uncertain or falls within a specific confidence range, involve human underwriters for final review. This can help catch errors that the model might make, particularly in borderline cases.

**5. Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone. Comment on this.**

#### **Comment on Avoiding Risky Loan Disbursements Due to NPA Concerns**

Given the critical issue of Non-Performing Assets (NPA) in the lending industry, it is crucial to adopt a conservative approach when disbursing loans. NPAs can significantly impact a financial

institution's profitability and liquidity, as they represent loans that borrowers are unlikely to repay. This makes it imperative to minimize exposure to high-risk applicants.

### **Key Considerations:**

#### **1. Prioritize Conservative Risk Management:**

- Given the financial implications of NPAs, the lending strategy should prioritize minimizing risk over maximizing loan disbursements. The cost of approving a high-risk applicant can outweigh the potential earnings from interest if the loan turns into an NPA.

#### **2. Rigorous Applicant Screening:**

- Implementing stringent criteria for applicant evaluation is essential. By thoroughly analyzing applicants' credit histories, income stability, and other financial behaviors, the model can help identify and filter out high-risk applicants more effectively.

#### **3. Focus on Long-Term Stability:**

- Disbursing loans to applicants with a proven track record of financial responsibility, even if it means rejecting a larger number of applications, will help ensure long-term financial stability. This approach helps prevent short-term gains from turning into long-term losses due to NPAs.

#### **4. Iterative Model Improvement:**

- Continuously refine the predictive model based on emerging trends in borrower behavior and economic conditions. This will allow for better anticipation of potential defaults and reduce the likelihood of disbursing loans that could become NPAs.

#### **5. Balance Caution with Opportunity:**

- While it is crucial to avoid disbursing loans to high-risk individuals, it's also important not to be overly conservative. A balanced approach that includes regular updates to the model and consideration of new data sources can help identify safe lending opportunities without excessive risk.

#### **6. Regular Portfolio Monitoring:**

- Continuously monitor the loan portfolio to identify early signs of potential defaults. Early intervention strategies can be developed to mitigate losses, thereby preventing loans from becoming NPAs.

**Conclusion:** In light of the NPA problem, playing it safe by being selective in loan disbursements is a prudent strategy. This approach helps protect the financial institution from potential losses while ensuring that the loans granted are more likely to be repaid.

## **Recommendations**

Based on the analysis and model insights, the following recommendations can be made to improve loan disbursement practices while minimizing risk:

#### **1. Focus on High-Grade Applicants:**

- Prioritize lending to applicants with Grade A, B, and C ratings, as they have shown a higher likelihood of fully repaying their loans. This will reduce the overall risk in the loan portfolio.

## 2. **Prefer Shorter Loan Tenures:**

- Encourage applicants to choose 36-month loan terms, as these loans have demonstrated a higher success rate compared to 60-month loans. Shorter tenures reduce the likelihood of default, ensuring better cash flow and lower risk.

## 3. **Enhance Creditworthiness Checks:**

- Give more weight to applicants with stable employment, no derogatory remarks, and no history of bankruptcies. These factors are strong indicators of an applicant's ability to repay the loan, reducing the chances of default.

## 4. **Target Debt Consolidation Loans:**

- Focus on applicants seeking loans for debt consolidation, as this purpose often reflects a proactive approach to managing finances. Such loans are more likely to be repaid in full, contributing to a healthier loan portfolio.

## 5. **Consider Zip-Code Risk Factors:**

- Implement stricter lending criteria for applicants from high-risk zip codes with a history of defaults. Conversely, offer more favorable terms to applicants from low-risk areas where the likelihood of loan repayment is higher.

## 6. **Regularly Update the Model:**

- Continuously retrain and update the logistic regression model with new data to reflect changes in economic conditions and borrower behavior. Regular updates will ensure the model remains accurate and relevant.

## 7. **Monitor Key Features:**

- Pay close attention to the most important features identified by the model, such as zip code, grade, debt-to-income ratio (DTI), and loan amount. These features should be regularly monitored to refine loan approval criteria.

## 8. **Consider Implementing a Two-Stage Model:**

- Utilize a two-stage model where the first stage identifies potential defaulters with high recall, and the second stage focuses on reducing false positives. This approach balances the need for accurate default detection with the goal of minimizing missed opportunities for lending.

## 9. **Balance Loan Approvals with NPA Risks:**

- While aiming to reduce NPAs, it is also important not to be overly conservative in loan disbursements. Striking the right balance will ensure that opportunities for profitable lending are not missed.

## 10. **Implement Human Review for High-Risk Cases:**

- For applicants who are borderline in terms of risk, consider involving human underwriters to review the cases. This additional step can help prevent potential NPAs while still allowing for reasonable loan approvals.

## Questionnaire:

1. What percentage of customers have fully paid their Loan Amount?

- **80%**

2. Comment about the correlation between Loan Amount and Installment features.

- **High Co-relation: 0.97**

3. The majority of people have home ownership as \_\_\_\_.

- **MORTGAGE**

4. People with grades 'A' are more likely to fully pay their loan. (T/F)

- **True.** *Although grade B has highest applicants but grade A high loan payback rate.*

5. Name the top 2 afforded job titles.

- **Teacher and Manager**

6. Thinking from a bank's perspective, which metric should our primary focus be on (ROC AUC, Precision, Recall, F1 Score)

- In my opinion **Recall**. Because higher number of defaulters(False Negative) cause high financial loss.
- But if Loan Tap is financially strong and wants to gain more interest than **F1-score** is also a good choice.

7. How does the gap in precision and recall affect the bank?

- High Recall and Low Precision indicates lower number of Non-Performing Assest at the cost of gaining more interest.
- Low Recall and High Precision indicates chance of gaining more interest at the cost of high number of Non-Performing Assest.

8. Which were the features that heavily affected the outcome?

- **Address (Zip codes), Grade, DTI (debt-to-income ratio) and loan amount**
- **Installment** is also a feature that heavily affected the outcome as it highly depends on Loan Amount.

9. Will the results be affected by geographical location? (Yes/No)

- **Yes**
- Zip-code: **05113, 00813, 29597** are the areas where almost every applicant has paid their loan successfully.
- Zip-code **11650, 86630, 93700** are the areas where almost every applicant has defaulted.